

PROJECT CONCEPT · CONFIDENTIAL DRAFT

A Support Tool for Understanding Intrusive Thoughts & Compulsive Patterns

Helping a person notice what happens in the moment, and understand it later — while they remain the one who draws the conclusions.

Foundational concept document

Scope: purpose, users, problem, features, and the reasoning behind each core design choice

Status: early-stage · not a clinical or engineering specification

1 What this is

A private application that helps a person observe and make sense of their own intrusive thoughts and compulsive patterns.

Intrusive thoughts and compulsions often begin automatically. From the inside they feel like a single, unavoidable event rather than a chain of connected parts, which makes them hard to understand and hard to respond to differently. This application exists to make that inner process visible. It gives the person a place to **capture** what is happening in the moment, and later to **understand** it — noticing how thoughts, emotions, triggers, past experiences, and responses connect.

The application is deliberately a **support tool, not a diagnostic one**. It does not tell anyone that they have a disorder, and it does not treat the person as a patient. It provides structure, memory, questions, and visualisation; the person does the work of understanding themselves. The tool is available, never insistent.

GUIDING STANCE

The application does not do the work for the person, and it does not pressure them to return. There are no streaks, no guilt, no engagement tricks. If a person chooses not to use it, that is their choice. The app offers structure and stays out of the way.

2 Who it is for

The application is for **anyone who experiences OCD-related patterns — whether or not they have a formal diagnosis** — and who believes a tool like this could help them understand or manage those patterns. A clinical label is not required to feel that the app is for them; what matters is that the person recognises the pattern in themselves.

Because it does not diagnose, the application never sorts people into categories such as “patient” or “sufferer.” It speaks in ordinary, non-clinical language — about *patterns, experiences, and what the person is noticing* — rather than about disorders or symptoms.

3 The problem it solves

Intrusive thoughts and compulsions run automatically and fast. The person rarely gets a clear look at the machinery: a trigger arrives, a thought follows, an emotion charges it, and a compulsion offers the only apparent exit — all before there is any room to reflect. This opacity is the core problem, because a person cannot respond thoughtfully to something they cannot see.

The application addresses this by helping the person notice, record, and connect the parts — thoughts, emotions, triggers, past experiences, and responses — so that what felt like one automatic surge becomes a chain with relationships that can be examined. The intended result is twofold: **a clearer view of one’s own patterns**, and **a changed response** — less confusion in the moment, and less pull toward the compulsion once it is seen as one option in a chain rather than the only way out.

4 How the person uses it: two modes

The application recognises that the same person has very different capacities depending on their state, so it offers two distinct modes rather than forcing one interface to serve both.

Capture now — fast, low-effort, no analysis required

During a difficult moment the person has little spare attention, so the app gets out of the way. In a few taps, a voice note, or one short sentence, they can record the thought, emotion, trigger, and urge without being asked to analyse anything. Fields can be left blank; capture tolerates incompleteness. This quick entry is not disposable — it is the seed that later reflection grows from.

Understand later — slower, reflective, when the person is calmer

When there is mental room, the app invites deeper reflection: what happened, what the person did, what reduced or reinforced the pattern, and whether it connects to earlier experiences. Here the app can be curious and patient. Crucially, **the person draws the connections themselves**; the app surfaces material and asks questions, but it does not hand over conclusions.

WHY THIS SPLIT MATTERS

Capturing a moment while a compulsion is active is a fundamentally different task from reflecting on it when calm. Designing them as one experience would compromise both. This distinction is the spine of the whole design.

5 Features

Feature	What the person can do
Capture Now	Quickly type or record what is happening — thought, emotion, trigger, urge — without needing to analyse it. Built for speed and minimal effort during a spike.
Reflective chatbot	Talk more openly when calmer. The chatbot does not simply comfort or agree; it asks questions, identifies unclear points, and helps the person examine what happened and draw their own connections.
Pattern graph	A growing map of thoughts, emotions, triggers, behaviours, and past experiences. The system may suggest connections or categories, but the person confirms, rejects, edits, rates, and explores them.
Personalised memory	The app remembers the person's recurring patterns and the way they communicate, so reflection over time builds a progressively truer picture of themselves rather than starting fresh each time.
Meditation & reflection space	A deliberately minimal space: a timer and an area for notes. The app does not decide what calm should look or sound like — the person may choose silence, calm music, a favourite artist, or anything else.
Turn-toward-help safety net	In rare, serious situations the app gently points the person toward appropriate human help, rather than acting as an authority or a reporter.

6 The graph, and why we use it

The graph is the visible form of the word *connect*. It represents the person's inner material as nodes — thoughts, emotions, triggers, behaviours, past experiences — with links between them. It is built **progressively, together with the person**, not generated for them in one step.

We use a graph because it does something plain text cannot: it turns a scattered history of moments into a structure the person can actually see and navigate. Patterns that are invisible inside a single entry become apparent when the same trigger, or the same internal process, connects across many entries.

The person remains the authority

The system may *suggest* a connection or a category, but a suggestion arrives as something tentative — a proposal to be made real, not a verdict to be deleted. The person confirms, rejects, edits, rates, and explores. This keeps the insight with the person: the moment the app declares the pattern for them, they stop doing the connecting, and the connecting is the point.

BEYOND TOPIC-MATCHING

The graph's most valuable job is to surface a **shared internal process across situations that look unrelated**. A conflict at work and a conflict with a parent may appear different on the surface yet involve the same inner machinery — for example uncertainty, self-blame, repeated analysis, reassurance-seeking, or avoidance. The app can gently raise this deeper similarity as a question; the person confirms, rejects, or corrects it.

7 Why we use a large language model

A large language model lets the reflective mode respond to what the person *actually* wrote, rather than offering fixed, generic prompts. It can notice that today's entry echoes one from weeks ago, ask the question that fits this specific moment, and reflect a pattern back in the person's own words. That responsiveness is the natural payoff of building on a model rather than a static form.

But the model is shaped deliberately against its own defaults, because two of them would undermine the app:

- **It must question, not conclude.** A model's instinct is to be helpful by summarising and deciding. Here that would steal the person's insight. The reflection is mostly interrogative — “you noted work stress here too; do you see a connection?” — leaving the person to draw the line.
- **It must not become a source of reassurance.** Reassurance-seeking is itself a compulsion. A chatbot that reflexively soothes (“don't worry, that thought means nothing”) can become a target for the very ritual the app is meant to loosen. Sometimes the most helpful move is to acknowledge the thought without litigating whether it is “true,” and without giving the reassurance hit.

DESIGN RISK TO HOLD

The chatbot and the graph could otherwise *collude* into a compulsion — the same node re-opened and the same question re-asked many times a day as a checking ritual through a new surface. The system should be able to gently notice “we have returned here often” rather than dutifully re-examining on demand.

8 Personalisation and fine-tuning

“Personalisation” here means the app adapts to how the individual communicates and to their recurring patterns. It is important to be precise about how this is achieved, because the obvious approach is the wrong one.

Per-person adaptation is memory, not model training

Adapting to a specific person is done by **remembering and retrieving their material** — their confirmed graph, their recurring process, their vocabulary and rhythm — and bringing the relevant pieces into the conversation when needed. The underlying model itself is not retrained on any individual. This is what allows the app to feel personal while keeping the person’s data controllable and deletable, and while ensuring the provider of the model never keeps or learns from it.

MEMORY IS PROCESS-FIRST, NOT CONTENT-FIRST

What the app remembers most durably is the **pattern and process** (for example “uncertainty → self-blame → checking”), not the exact frightening wording of an intrusive thought. This is better for privacy and better for the person: re-surfacing distressing raw text can feed the very loop the app is trying to loosen. The rule: adapt to how the person talks; never absorb the obsessional logic itself.

Where fine-tuning may belong

Fine-tuning, if used, is reserved for shaping the app’s **overall clinical character** — **shared by everyone and containing no individual’s private material**: holding an unwanted thought without deciding whether it is true, declining the reassurance reflex, asking rather than declaring, using non-clinical language, and turning the person toward help when needed. Even this is something to approach later and carefully, only once there is a well-reviewed body of examples; a strong, well-designed set of instructions and examples comes first.

9 Why the app turns the person toward help

The application deals with obsessive-compulsive patterns, and it must be extremely careful about one thing above all: **an intrusive thought is not an intention**. A person may say that they have an intrusive thought about causing terrible harm. In OCD, such a thought is typically *unwanted* and *distressing* — the opposite of what the person wants — and the distress itself is the signal that it is an obsession, not a plan. It does not mean anything happened, or that anything will.

Because of this, the app is built so that **no one is ever framed, reported, or made to feel in danger for disclosing a thought**. The entire tool depends on it being safe to speak the unwanted thought aloud; OCD feeds on secrecy and avoidance, so a tool that felt like it might turn a person in would fail at its core purpose.

HOW THE SAFETY NET BEHAVES

In the rare situation that genuinely warrants it — recognised by signs of real intent and planning, never by the content of a thought alone, and only at a deliberately high threshold — the app **gently turns the person toward appropriate human help** (such as a crisis line or a trusted contact). It offers a resource; it does not act as an authority and does not report the person. This safety response is kept separate from personalisation, so it behaves consistently for everyone rather than adapting to any individual.

10 Feedback, rating, templates, and questions

The person is an active participant, so the app gives them clear ways to steer it and correct it.

Feedback and rating

- **On graph suggestions:** confirm, reject, edit, rate, or explore any connection or category the system proposes. Rejection and correction are first-class actions, not afterthoughts.
- **On reflections and questions:** the person can indicate when a question or observation was useful or off the mark, so the experience improves and stays in their own terms.

Templates and questions

Two kinds of scaffolding support the two modes:

- **Capture templates** — the lightest possible structure for a difficult moment: a short field for the thought, and optional quick selections for emotion, trigger, and urge. Nothing is mandatory.
- **Reflection questions** — used in the calmer mode to help the person examine an entry: what happened, what they did, what reduced or reinforced the pattern, and whether it connects to earlier experiences. These are open and interrogative, chosen to fit the specific entry rather than applied uniformly.

11 What data we need, and how it is useful

The application works with a small set of information types, each tied directly to a purpose. The principle throughout is **data minimisation**: never hold more than the task requires.

Information	Why it is useful
In-the-moment captures	Thought, emotion, trigger, urge. The raw material that later reflection grows from, and the basis for seeing patterns over time.
Reflections	What happened, what the person did, what helped or worsened it. Turns single moments into understanding and feeds the graph.
Confirmed graph	The person's own map of connections. The durable record of their self-understanding, and what makes personalisation meaningful.
Recurring process	The internal patterns (e.g. uncertainty, self-blame, checking) held in preference to distressing raw wording. Powers cross-situation insight while protecting the person.
Communication style	Vocabulary and rhythm, so reflection sounds like the person and not a generic script.
Feedback signals	Confirmations, rejections, ratings. Keeps the app accurate and keeps the person in authority over it.

DATA PROTECTION AS A CORE REQUIREMENT

This is among the most sensitive information a person can hold. Consent (confirmed before sign-up), user control, deletion, security, and legal compliance — to be refined against European Union law — are treated as foundational design requirements, not later additions. Where a third-party model is used, it is arranged so that the person's content is never stored by it and never used to train it.

Specific data-handling architecture, provider terms, the clinical review of the app's reflective stance, and the full text of consent and terms are deliberately out of scope for this concept document and are to be defined in a later stage.

12 In one paragraph

This is a private, non-diagnostic support tool for anyone who experiences OCD-related patterns. It lets a person **capture** intrusive thoughts and compulsions quickly in the moment, and **understand** them later through reflection, a questioning chatbot, and a graph of their thoughts, emotions, triggers, behaviours, and past experiences that they build and control themselves. A language model makes the reflection responsive and personal — adapting through memory rather than by training on the person — while never concluding for them and never reassuring them into a new compulsion. When something genuinely serious appears, judged by intent rather than by the content of a thought, the app gently turns the person toward human help without ever framing or reporting them. Throughout, the person does the understanding; the app provides structure, memory, questions, and visualisation, and protects the person's data as its first responsibility.